

VIETNAM NATIONAL UNIVERSITY, HO CHI MINH CITY
UNIVERSITY OF TECHNOLOGY
FACULTY OF COMPUTER SCIENCE AND ENGINEERING



OPERATING SYSTEMS

Report for Assignment #02 Simple Operating System

Advisor: Tran Viet Toan
Students: Vo Minh Khoa - 1812670
Tran Long Vi - 1814804

HO CHI MINH CITY, DECEMBER 2020



Contents

1	Scheduling	2
1.1	Questions	2
1.2	Implementation	3
1.2.1	Priority queue	3
1.2.2	Scheduler	4
1.3	Gantt Diagrams	7
1.3.1	Test 1	7
1.3.2	Test 2	7
2	Memory management	8
2.1	Questions	8
2.2	Implement and show the status of RAM	8
2.3	Explain file input m0	12
3	Overall	12
4	Simulation	15

1 Scheduling

1.1 Questions

Question 1: What is the advantage of using priority feedback queue in comparison with other scheduling algorithms you have learned?

The priority feedback queue algorithm: The Priority feedback queue is based on the multilevel feedback queue algorithm used on the Linux kernel.

The multilevel feedback queue (MFQ) algorithm: In the Multilevel Feedback Queue system, the scheduler can move processes between queues according to its observed properties, changing the process's priority.

we can apply MFQ to the priority feedback queue algorithm that applied MFQ by using two queues with priority: ready queue and run queue.

- Ready queue is a priority queue that inherits a portion of Priority Scheduling (PS), every time the CPU receives the process with the highest priority from the ready queue.
- Run queue contains processes that are waiting to continue to execute after their slots have expired but their process has not yet been completed.

Ready queue has a higher priority than the run queue so it is executed first by the CPU. When the CPU moves to the next slot, it looks for the process in the ready queue.

When the ready queue is empty, the processes on the run queue switch to the ready queue to consider the next slot. In this situation, we promote Process.

When the process at the ready queue runs out of the given quantum time it goes down to the run queue. Both queues are the priority queue, the priority level is based on the priority level of the process in the queue. In this case, we downgrade Process.

Disadvantages of other schedulings are:

- First Come First Served (FCFS):
For this algorithm, due to using the nonpreemptive strategy in the process, if a process starts, the CPU executes the process until it finishes. Therefore, the processes behind (possibly the processes with a short CPU burst) in the queue have to wait longer.
- Shortest Job First (SJF):
 - + It is necessary to estimate the amount of time it takes for the next process of the CPU to first while this is often not possible in practice.
 - + Process with long CPU burst has more waiting time or indefinitely delay when there are many processes with short CPU burst in queue.
- Round Robin (RR):
 - + Throughput is often large depending on the time quantum. If time quantum is too large, RR will like FCFS.
 - + If the time quantum is too small, the CPU context switch will increase a lot causing OS overhead and reduce CPU utilization.
- Priority Scheduling (PS):
 - + There must be a scheduler with processes with equal priorities.
 - + Processes with lower priorities may not have a chance to be executed.

The priority feedback queue (PFQ) algorithm can overcome the disadvantages of other scheduling with the following points:

- This algorithm is based on the Multilevel Feedback Queue, which uses two different queues and moves the processes back and forth between queues until the process is completed. So respond time can be shorter.
- Turnaround time is optimized: PFQ runs the process according to the time quantum then changes the process's priority, so it learns from the process's past behavior and then predicts its future behavior. In this way, the PFQ tries to run a short process first, so to optimize the turnaround time.
- It's also more flexible than the scheduling algorithms have been learned.
- Ensure fairness for processes by applying Round Robin style.
- Avoid starvation: Since PFQ also needs to prioritize processes with high priority, if there are no different queues, the process with low priority will not take its turn, it will still be executed before the processes with higher priority after the slot has been completed.

1.2 Implementation

1.2.1 Priority queue

We use *enqueue()* and *dequeue()* in *queue.c* to implement priority queue in this project.

- *enqueue()*: Put a new process into queue if queue is empty.

```
9  void enqueue(struct queue_t * q, struct pcb_t * proc) {
10      /* TODO: put a new process to queue [q] */
11      if(q->size < MAX_QUEUE_SIZE)
12      {
13          q->proc[q->size] = proc;
14          q->size++;
15      }
16  }
```

Figure 1: *enqueue()* implementation

- *dequeue()*: Retrieve process with highest priority in queue, update queue status when retrieving element.

```
18 struct pcb_t * dequeue(struct queue_t * q) {
19     if(q->size!=0)
20     {
21         struct pcb_t* temp_proc=NULL;
22         int max_priority=0;
23         int max_idx=0;
24         int i;
25         for(i=0;i<q->size;++i)
26         {
27             if(q->proc[i]->priority>max_priority)
28             {
29                 temp_proc=q->proc[i];
30                 max_priority=temp_proc->priority;
31                 max_idx=i;
32             }
33         }
34         for(i=max_idx;i<q->size-1;++i)
35             q->proc[i]=q->proc[i+1];
36         q->proc[q->size-1]=NULL;
37         --q->size;
38         return temp_proc;
39     }
40     return NULL;
41 }
```

Figure 2: *dequeue()* implementation

1.2.2 Scheduler

The scheduler is used to manage updating the processes that will be executed for the CPU.

We use *get_proc()* in *sched.c* to implement scheduler in this project.

get_proc() will return the process at queue “ready”. If queue is empty at the time function is called, update queue again with processes that are waiting for next slot in queue “run”. On the contrary, we find the process with high priority from this queue.

```
20 struct pcb_t * get_proc(void) {
21     struct pcb_t * proc = NULL;
22     pthread_mutex_lock(&queue_lock);
23     if(ready_queue.size==0)
24     {
25         int i = 0
26         while(i <= MAX_QUEUE_SIZE)
27         {
28             i++;
29             ready_queue.proc[i]=run_queue.proc[i];
30             run_queue.proc[i]=NULL;
31         }
32         ready_queue.size=run_queue.size;
33         run_queue.size=0;
34     }
35     proc=dequeue(&ready_queue);
36     pthread_mutex_unlock(&queue_lock);
37     return proc;
38 }
```

Figure 3: *get_proc()* implementation

After run command *make test_sched* by terminal, we can see result like this:

```
----- SCHEDULING TEST 0 -----
./os sched_0
Time slot 0
    Loaded a process at input/proc/s0, PID: 1
Time slot 1
    CPU 0: Dispatched process 1
Time slot 2
Time slot 3
    CPU 0: Put process 1 to run queue
    CPU 0: Dispatched process 1
Time slot 4
    Loaded a process at input/proc/s1, PID: 2
Time slot 5
    CPU 0: Put process 1 to run queue
    CPU 0: Dispatched process 2
Time slot 6
Time slot 7
    CPU 0: Put process 2 to run queue
    CPU 0: Dispatched process 2
Time slot 8
Time slot 9
    CPU 0: Put process 2 to run queue
    CPU 0: Dispatched process 1
Time slot 10
Time slot 11
    CPU 0: Put process 1 to run queue
    CPU 0: Dispatched process 2
Time slot 12
Time slot 13
    CPU 0: Put process 2 to run queue
    CPU 0: Dispatched process 1
Time slot 14
Time slot 15
    CPU 0: Put process 1 to run queue
    CPU 0: Dispatched process 2
Time slot 16
    CPU 0: Processed 2 has finished
    CPU 0: Dispatched process 1
Time slot 17
Time slot 18
    CPU 0: Put process 1 to run queue
    CPU 0: Dispatched process 1
Time slot 19
Time slot 20
    CPU 0: Put process 1 to run queue
    CPU 0: Dispatched process 1
Time slot 21
Time slot 22
    CPU 0: Put process 1 to run queue
    CPU 0: Dispatched process 1
Time slot 23
    CPU 0: Processed 1 has finished
    CPU 0 stopped
```

Figure 4: *Scheduling Test 1 result*

```
----- SCHEDULING TEST 1 -----
./os sched 1
Time slot 0
    Loaded a process at input/proc/s0, PID: 1
Time slot 1
    CPU 0: Dispatched process 1
Time slot 2
Time slot 3
    CPU 0: Put process 1 to run queue
    CPU 0: Dispatched process 1
Time slot 4
    Loaded a process at input/proc/s1, PID: 2
Time slot 5
    CPU 0: Put process 1 to run queue
    CPU 0: Dispatched process 2
Time slot 6
    Loaded a process at input/proc/s2, PID: 3
Time slot 7
    CPU 0: Put process 2 to run queue
    CPU 0: Dispatched process 3
    Loaded a process at input/proc/s3, PID: 4
Time slot 8
Time slot 9
    CPU 0: Put process 3 to run queue
    CPU 0: Dispatched process 4
Time slot 10
Time slot 11
    CPU 0: Put process 4 to run queue
    CPU 0: Dispatched process 2
Time slot 12
Time slot 13
    CPU 0: Put process 2 to run queue
    CPU 0: Dispatched process 3
Time slot 14
Time slot 15
    CPU 0: Put process 3 to run queue
    CPU 0: Dispatched process 1
Time slot 16
Time slot 17
    CPU 0: Put process 1 to run queue
    CPU 0: Dispatched process 4
Time slot 18
Time slot 19
    CPU 0: Put process 4 to run queue
    CPU 0: Dispatched process 2
Time slot 20
Time slot 21
    CPU 0: Put process 2 to run queue
    CPU 0: Dispatched process 3
Time slot 22
```

Figure 5: *Scheduling Test 2 result (1)*

```
Time slot 24
Time slot 25
    CPU 0: Put process 1 to run queue
    CPU 0: Dispatched process 4
Time slot 26
Time slot 27
    CPU 0: Put process 4 to run queue
    CPU 0: Dispatched process 2
Time slot 28
    CPU 0: Processed 2 has finished
    CPU 0: Dispatched process 3
Time slot 29
Time slot 30
    CPU 0: Put process 3 to run queue
    CPU 0: Dispatched process 1
Time slot 31
Time slot 32
    CPU 0: Put process 1 to run queue
    CPU 0: Dispatched process 4
Time slot 33
Time slot 34
    CPU 0: Put process 4 to run queue
    CPU 0: Dispatched process 3
```

Figure 6: *Scheduling Test 2 result (2)*

```

Time slot 35
Time slot 36
    CPU 0: Put process 3 to run queue
    CPU 0: Dispatched process 1
Time slot 37
Time slot 38
    CPU 0: Put process 1 to run queue
    CPU 0: Dispatched process 4
Time slot 39
Time slot 40
    CPU 0: Put process 4 to run queue
    CPU 0: Dispatched process 3
Time slot 41
Time slot 42
    CPU 0: Processed 3 has finished
    CPU 0: Dispatched process 1
Time slot 43
Time slot 44
    CPU 0: Put process 1 to run queue
    CPU 0: Dispatched process 4
Time slot 45
    CPU 0: Processed 4 has finished
    CPU 0: Dispatched process 1
Time slot 46
    CPU 0: Processed 1 has finished
    CPU 0 stopped
  
```

Figure 7: Scheduling Test 2 result (3)

1.3 Gantt Diagrams

In this Assignment, my team draw Gantt Diagrams to describe how processes are executed by the CPU in 2 situations.

1.3.1 Test 1

In this situation, CPU executes 2 processes p1 and p2 in 23 time slots.

Gantt Diagrams:





Figure 8: Gantt diagram in test 1

1.3.2 Test 2

In this situation, CPU executes 4 processes p1, p2, p3 and p4 in 48 time slots.

Gantt Diagrams:



2 Memory management

2.1 Questions

Question 2: What is the advantage and disadvantage of segmentation with paging?
Segmentation with paging is a mechanism Virtual Memory Engine (VME) used to manage memory.

The segmentation technique satisfies the program's need to demonstrate the logical structure of the program, but it leads to the situation of having to allocate memory blocks of different sizes for segments in the physical memory. Therefore, if we do Paging the segmentation, the problem will be better resolved.

Segmentation with paging:

- Address space is a set of segments, each segment is divided into many pages.
- When a process is put into OS, it will allocate to the process necessary pages to contain all segments of the process.
- Cuz the process uses virtual addresses to access RAM, we should set the virtual addresses of the allocated pages to be adjacent.

Advantage of Segmentation with paging:

- Saving memory, using memory effectively.
- Allocating intermittent memory in a simpler way.
- Taking advantages of Paging and Segmentation and minimize disadvantages of each other.
- Sharing data between processes is more flexible
- Fixing the size of the page too large by paging in each segment.
- There is no foreign fragmentation.

Disadvantage of Segmentation with paging:

- Size of the process is limited by the size of physical memory.
- It is difficult to maintain multiple processes at the same time in memory.
- Page tables need a lot of memory space, so it is not good for systems with small RAM.

2.2 Implement and show the status of RAM

We use *alloc_mem()* and *free_mem()* in *mem.c* to implement the status of RAM in this project.
alloc_mem(): Allocation memory for the process and save the address of the first byte in the allocated memory region.

```
103 addr_t alloc_mem(uint32_t size, struct pcb_t * proc) {
104     pthread_mutex_lock(&mem_lock);
105     addr_t ret_mem = 0;
106     uint32_t num_pages = ((size % PAGE_SIZE) == 0) ? size / PAGE_SIZE :
107         size / PAGE_SIZE + 1;
108     int mem_avail = 0;
109     int i;
110     int num_avail_pages = 0;
111     for(i = 0; i < NUM_PAGES; i++){
112         if(_mem_stat[i].proc == 0){
113             num_avail_pages++;
114             if(num_avail_pages == num_pages && proc->bp + num_pages * PAGE_SIZE <= RAM_SIZE){
115                 mem_avail = 1;
116                 break;
117             }
118         }
119     }
120 }
121
122 if (mem_avail) {
123     ret_mem = proc->bp;
124     proc->bp += num_pages * PAGE_SIZE; int num_alloc_pages = 0;
125     int pre_index;
126     addr_t cur_vir_addr;
127     int seg_idx, page_idx;
128     for(i = 0; i < NUM_PAGES; i++){
129         if(_mem_stat[i].proc == 0){
130             _mem_stat[i].proc = proc->pid;
131             _mem_stat[i].index = num_alloc_pages;
132
133             if(_mem_stat[i].index != 0)
134                 _mem_stat[pre_index].next = i;
135             pre_index = i;
136         }
137     }
138 }
```

Figure 9: *alloc_mem()* implementation(1)

```
126     int found = 0;
127     struct seg_table_t * seg_table = proc->seg_table;
128     if(seg_table->table[0].pages == NULL)
129         seg_table->size = 0;
130
131     cur_vir_addr = ret_mem + (num_alloc_pages << OFFSET_LEN) ;
132
133     seg_idx = get_first_lv(cur_vir_addr);
134     page_idx = get_second_lv(cur_vir_addr);
135     int j;
136     for(j = 0; j < seg_table->size; j++){
137         if(seg_table->table[j].v_index == seg_idx){
138             struct page_table_t * cur_page_table = seg_table->table[j].pages;
139
140             cur_page_table->table[cur_page_table->size].v_index = page_idx;
141             cur_page_table->table[cur_page_table->size].p_index = i;
142
143             cur_page_table->size++;
144
145             found = 1;
146             break;
147         }
148     }
149
150     if(!found){
151         seg_table->table[seg_table->size].v_index = seg_idx;
152         seg_table->table[seg_table->size].pages = (struct page_table_t *)malloc(sizeof
153
154         seg_table->table[seg_table->size].pages->table[0].v_index = page_idx;
155         seg_table->table[seg_table->size].pages->table[0].p_index = i;
156
157         seg_table->table[seg_table->size].pages->size = 1;
158
159         seg_table->size++;
160     }
```

Figure 10: *alloc_mem()* implementation(2)

```

162     num_alloc_pages++;
163     if(num_alloc_pages == num_pages){
164         _mem_stat[i].next = -1;
165         break;
166     }
167 }
168 }
169 }
170
171 pthread_mutex_unlock(&mem_lock);
172 return ret_mem;
173 }

```

Figure 11: *alloc_mem()* implementation(3)

free_mem(): Release memory region allocated.

```

175 int free_mem(addr_t address, struct pcb_t * proc) {
176     pthread_mutex_lock(&mem_lock);
177
178     struct page_table_t * page_table = get_page_table(get_first_lv(address), proc->seg_table);
179
180     int valid = 0;
181     if(page_table != NULL){
182         int i;
183         for(i = 0; i < page_table->size; i++){
184             if(page_table->table[i].v_index == get_second_lv(address)){
185                 addr_t physical_addr;
186                 if(translate(address, &physical_addr, proc)){
187                     int p_index = physical_addr >> OFFSET_LEN;
188                     int num_free_pages = 0;
189                     addr_t cur_vir_addr = (num_free_pages << OFFSET_LEN) + address;
190                     addr_t seg_idx, page_idx;
191                     do{
192                         _mem_stat[p_index].proc = 0;
193                         int found = 0;
194                         int k;
195                         seg_idx = get_first_lv(cur_vir_addr);
196                         page_idx = get_second_lv(cur_vir_addr);
197                         for(k = 0; k < proc->seg_table->size && !found; k++){
198                             if(proc->seg_table->table[k].v_index == seg_idx){
199                                 int l;
200                                 for(l = 0; l < proc->seg_table->table[k].pages->size; l++){
201                                     if(proc->seg_table->table[k].pages->table[l].v_index == page_idx){
202                                         int m;
203                                         for(m = 1; m < proc->seg_table->table[k].pages->size - 1;
204                                             proc->seg_table->table[k].pages->table[m] = proc->seg_
205
206                                     proc->seg_table->table[k].pages->size--;

```

Figure 12: *free_mem()* implementation(1)

```

207     if(proc->seg_table->table[k].pages->size == 0){
208         free(proc->seg_table->table[k].pages);
209         for(m = k; m < proc->seg_table->size - 1; m++){
210             proc->seg_table->table[m] = proc->seg_table->table
211             proc->seg_table->size--;
212         }
213         found = 1;
214         break;
215     }
216 }
217 }
218 }
219 }
220     p_index = _mem_stat[p_index].next;
221     num_free_pages++;
222 }
223 while(p_index != -1);
224 valid = 1;
225 }
226 break;
227 }
228 }
229 }
230 pthread_mutex_unlock(&mem_lock);
231
232 if(!valid)
233     return 1;
234 else
235     return 0;
236 }
237 }

```

Figure 13: *free_mem()* implementation(2)

After run command *make test_mem* by terminal, we can see result like this:

```

khoa@ubuntu:~/Desktop/Assignment2-05/source_code$ make mem
gcc -Iinclude -Wall -c -g src/paging.c -o obj/paging.o
gcc -Iinclude -Wall -g obj/paging.o obj/mem.o obj/cpu.o obj/loader.o -o mem -lpthread
khoa@ubuntu:~/Desktop/Assignment2-05/source_code$ make test_mem
----- MEMORY MANAGEMENT TEST 0 -----
./mem input/proc/m0
000: 00000-003ff - PID: 01 (idx 000, nxt: 001)
      003e8: 15
001: 00400-007ff - PID: 01 (idx 001, nxt: -01)
002: 00800-00bff - PID: 01 (idx 000, nxt: 003)
003: 00c00-00fff - PID: 01 (idx 001, nxt: 004)
004: 01000-013ff - PID: 01 (idx 002, nxt: 005)
005: 01400-017ff - PID: 01 (idx 003, nxt: 006)
006: 01800-01bff - PID: 01 (idx 004, nxt: -01)
014: 03800-03bff - PID: 01 (idx 000, nxt: 015)
      03814: 66
015: 03c00-03fff - PID: 01 (idx 001, nxt: -01)
NOTE: Read file output/m0 to verify your result
----- MEMORY MANAGEMENT TEST 1 -----
./mem input/proc/m1
NOTE: Read file output/m1 to verify your result (your implementation should print nothing)
khoa@ubuntu:~/Desktop/Assignment2-05/source_code$

```

Figure 14: *test_mem* result

2.3 Explain file input m0

```
1 7
alloc 13535 0
alloc 1568 1
free 0
alloc 1386 2
alloc 4564 4
write 102 1 20
write 21 2 1000
```

- In first line, “1” is the priority of process m0, and “7” is the instruction number of the input file.
- In 2nd line, comment “alloc 13535” will allocate 14 `_mem_stat` pages (from 000 to 013) and store the address of the first allocated byte in register 0.
- In 3rd line, comment “alloc 1568” will allocate 2 pages (014 and 015) and store the address of the first allocated byte in register number 1.
- In 4th line, comment “free 0” will release the allocated memory from the alloc command in register number 0, which means there are only pages 014 and 015 at this time.
- In 5th line, comment “alloc 1386” will allocate 2 pages (000 and 001), then store the address of the first allocated byte in register number 2.
- In 6th line, comment “alloc 4564” will allocate 5 pages (from 002 to 006) and store the address of the first allocated byte in register number 4.
- In 7th line, comment “write 102 1 20” will write the value 102 to the place where the address is equal to register address 1 plus an offset of 20 that will output 0x03814.
The physical address can be calculated by $Physical\ address = Base\ address + Offset$, in which, $Base\ address = 0x03800$ is the first byte and $offset = 20$.
- In final line, similarly, the result would be 003e8.

3 Overall

In file `mem.c`, in addition to the `alloc_mem()` and `free_mem()`, we also implement `get_page_table()` and `translate()`. Besides, we fixed `read_mem()` and `write_mem()`, too.

get_page_table() is used to find the paging table from segment:

```
47 static struct page_table_t *get_page_table(  
48     addr_t index,  
49     struct seg_table_t *seg_table)  
50 {  
51     int i;  
52     for (i = 0; i < seg_table->size; i++)  
53     {  
54         // Enter your code here  
55         if (index == seg_table->table[i].v_index)  
56         {  
57             return seg_table->table[i].pages;  
58         }  
59     }  
60     return NULL;  
61 }
```

Figure 15: *get_page_table()* implementation

translate() is used to map virtual addresses to physical addresses:

```
86 static int translate(  
87     addr_t virtual_addr, // Given virtual address  
88     addr_t *physical_addr, // Physical address to be returned  
89     struct pcb_t *proc)  
90 {  
91     addr_t offset = get_offset(virtual_addr);  
92     addr_t first_lv = get_first_lv(virtual_addr);  
93     addr_t second_lv = get_second_lv(virtual_addr);  
94     struct page_table_t *page_table = NULL;  
95     page_table = get_page_table(first_lv, proc->seg_table);  
96     if (page_table == NULL)  
97     {  
98         return 0;  
99     }  
100     int i;  
101     for (i = 0; i < 1 << PAGE_LEN; i++)  
102     {  
103         if (page_table->table[i].v_index == second_lv)  
104         {  
105             *physical_addr = page_table->table[i].p_index * PAGE_SIZE + offset;  
106             return 1;  
107         }  
108     }  
109     return 0;  
110 }
```

Figure 16: *translate()* implementation

read_mem() and *write_mem()* is added mutex to avoid asynchronous situation when they access to [ram] at the same time.

```
143 pthread_mutex_t mutex;
144 √ int read_mem(addr_t address, struct pcb_t * proc, BYTE * data) {
145     addr_t physical_addr;
146     √ if (translate(address, &physical_addr, proc)) {
147         pthread_mutex_lock(&mutex);
148         *data = _ram[physical_addr];
149         pthread_mutex_unlock(&mutex);
150         return 0;
151     }
152     √ else{
153         return 1;
154     }
155 }
156 √ int write_mem(addr_t address, struct pcb_t * proc, BYTE data) {
157     addr_t physical_addr;
158     √ if (translate(address, &physical_addr, proc)) {
159         pthread_mutex_lock(&mutex);
160         _ram[physical_addr] = data;
161         pthread_mutex_unlock(&mutex);
162         return 0;
163     }
164     √ else{
165         return 1;
166     }
167 }
```

Figure 17: *read_mem()* and *write_mem()* after added mutex

4 Simulation

After run command “*make test_all*” by Terminal, we can see result like below:

```
----- OS TEST 1 -----
./os os_1
Time slot 0
    Loaded a process at input/proc/p0, PID: 1
Time slot 1
    CPU 1: Dispatched process 1
Time slot 2
    Loaded a process at input/proc/s3, PID: 2
    CPU 3: Dispatched process 2
Time slot 3
    CPU 1: Put process 1 to run queue
    CPU 1: Dispatched process 1
    Loaded a process at input/proc/m1, PID: 3
Time slot 4
    CPU 0: Dispatched process 3
Time slot 5
    CPU 1: Put process 1 to run queue
    CPU 1: Dispatched process 1
    CPU 3: Put process 2 to run queue
    CPU 3: Dispatched process 2
    Loaded a process at input/proc/s2, PID: 4
Time slot 6
    CPU 0: Put process 3 to run queue
    CPU 0: Dispatched process 4
    CPU 3: Put process 2 to run queue
    CPU 3: Dispatched process 2
Time slot 7
    CPU 1: Put process 1 to run queue
    CPU 1: Dispatched process 3
    Loaded a process at input/proc/m0, PID: 5
    CPU 2: Dispatched process 1
Time slot 8
    CPU 0: Put process 4 to run queue
    CPU 0: Dispatched process 5
    Loaded a process at input/proc/p1, PID: 6
Time slot 9
    CPU 1: Put process 3 to run queue
    CPU 1: Dispatched process 6
    CPU 2: Put process 1 to run queue
    CPU 2: Dispatched process 4
    CPU 3: Put process 2 to run queue
    CPU 3: Dispatched process 3
Time slot 10
    CPU 0: Put process 5 to run queue
    CPU 0: Dispatched process 1
    Loaded a process at input/proc/s0, PID: 7
    CPU 2: Put process 4 to run queue
    CPU 2: Dispatched process 7
    CPU 3: Put process 3 to run queue
    CPU 3: Dispatched process 4
Time slot 11
    CPU 1: Put process 6 to run queue
    CPU 1: Dispatched process 2
```

Figure 18: *Final result (1)*


```
Time slot 12
    CPU 0: Processed 1 has finished
    CPU 0: Dispatched process 5
Time slot 13
    CPU 1: Put process 2 to run queue
    CPU 1: Dispatched process 3
    CPU 2: Put process 7 to run queue
    CPU 2: Dispatched process 7
    CPU 3: Put process 4 to run queue
    CPU 3: Dispatched process 2
Time slot 14
    CPU 0: Put process 5 to run queue
    CPU 0: Dispatched process 6
Time slot 15
    CPU 2: Put process 7 to run queue
    CPU 2: Dispatched process 4
    CPU 3: Put process 2 to run queue
    CPU 3: Dispatched process 7
    CPU 1: Processed 3 has finished
    CPU 1: Dispatched process 5
    Loaded a process at input/proc/s1, PID: 8
    CPU 0: Put process 6 to run queue
    CPU 0: Dispatched process 8
Time slot 16
    CPU 3: Put process 7 to run queue
    CPU 3: Dispatched process 7
Time slot 17
    CPU 2: Put process 4 to run queue
    CPU 2: Dispatched process 2
    CPU 1: Put process 5 to run queue
    CPU 1: Dispatched process 6
Time slot 18
    CPU 2: Processed 2 has finished
    CPU 2: Dispatched process 4
    CPU 0: Put process 8 to run queue
    CPU 0: Dispatched process 5
Time slot 19
    CPU 0: Processed 5 has finished
    CPU 0: Dispatched process 8
    CPU 3: Put process 7 to run queue
    CPU 3: Dispatched process 7
    CPU 1: Put process 6 to run queue
    CPU 1: Dispatched process 6
Time slot 20
    CPU 2: Put process 4 to run queue
    CPU 2: Dispatched process 4
```

Figure 19: *Final result (2)*

```
Time slot 21
  CPU 0: Put process 8 to run queue
  CPU 0: Dispatched process 8
  CPU 1: Put process 6 to run queue
  CPU 1: Dispatched process 6
  CPU 3: Put process 7 to run queue
  CPU 3: Dispatched process 7
Time slot 22
  CPU 2: Processed 4 has finished
  CPU 2 stopped
Time slot 23
  CPU 0: Put process 8 to run queue
  CPU 0: Dispatched process 8
  CPU 1: Processed 6 has finished
  CPU 1 stopped
  CPU 3: Put process 7 to run queue
  CPU 3: Dispatched process 7
Time slot 24
  CPU 0: Processed 8 has finished
  CPU 0 stopped
Time slot 25
  CPU 3: Put process 7 to run queue
  CPU 3: Dispatched process 7
Time slot 26
  CPU 3: Processed 7 has finished
  CPU 3 stopped
```

Figure 20: *Final result (3)*

```
MEMORY CONTENT:
000: 00000-003ff - PID: 05 (idx 000, nxt: 001)
      003e8: 15
001: 00400-007ff - PID: 05 (idx 001, nxt: -01)
002: 00800-00bff - PID: 05 (idx 000, nxt: 003)
003: 00c00-00fff - PID: 05 (idx 001, nxt: 004)
004: 01000-013ff - PID: 05 (idx 002, nxt: 005)
005: 01400-017ff - PID: 05 (idx 003, nxt: 006)
006: 01800-01bff - PID: 05 (idx 004, nxt: -01)
011: 02c00-02fff - PID: 06 (idx 000, nxt: 012)
012: 03000-033ff - PID: 06 (idx 001, nxt: 013)
013: 03400-037ff - PID: 06 (idx 002, nxt: 014)
014: 03800-03bff - PID: 06 (idx 003, nxt: -01)
021: 05400-057ff - PID: 01 (idx 000, nxt: -01)
      05414: 64
022: 05800-05bff - PID: 06 (idx 000, nxt: 023)
023: 05c00-05fff - PID: 06 (idx 001, nxt: 031)
024: 06000-063ff - PID: 05 (idx 000, nxt: 025)
      06014: 66
025: 06400-067ff - PID: 05 (idx 001, nxt: -01)
031: 07c00-07fff - PID: 06 (idx 002, nxt: 032)
      07de7: 0a
032: 08000-083ff - PID: 06 (idx 003, nxt: 033)
033: 08400-087ff - PID: 06 (idx 004, nxt: -01)
NOTE: Read file output/os_1 to verify your result
khoea@ubuntu:~/Desktop/Assignment2-OS/source_code$
```

Figure 21: *Final result (4)*